

Module 8: Integrations & Pipelines

In this module, you will:

1. **Learn about data integration patterns** with TimescaleDB
2. **Set up a realistic data ingestion pipeline** from REST APIs
3. **Export and analyze data with Python & Pandas**
4. **Create visualizations** with Matplotlib
5. **Build a complete data pipeline** from API to insights
6. **Clean up**

What are Integrations & Pipelines?

Simple Explanation

Data Integration = Getting data INTO TimescaleDB from external sources

Data Pipeline = The complete flow: Collect → Store → Process → Analyze → Visualize

Common Integration Patterns



Data Sources (Ingestion):

- REST APIs (weather, sensors, stock prices)
- File uploads (CSV, JSON)
- Direct database connections
- Scheduled data imports



Data Destinations (Export):

- Python/Pandas for analysis
- CSV/JSON exports for other tools
- Machine learning pipelines
- Real-time dashboards

Step 1: Project Setup

Create Project Directory

```
mkdir timescale-integrations  
cd timescale-integrations
```

Create docker-compose.yml

```

services:
  timescaledb:
    image: timescale/timescaledb:latest-pg16
    container_name: timescale-integrations
    environment:
      - POSTGRES_USER=admin
      - POSTGRES_PASSWORD=admin123
      - POSTGRES_DB=integrations_db
    ports:
      - "5432:5432"
    volumes:
      - timescale_data:/var/lib/postgresql/data
      - ./init-scripts:/docker-entrypoint-initdb.d

volumes:
  timescale_data:

```

Create Database Schema

Create `init-scripts/01-setup.sql`:

```

-- Enable TimescaleDB extension
CREATE EXTENSION IF NOT EXISTS timescaledb;

-- Weather data table
CREATE TABLE weather_data (
  time TIMESTAMPTZ NOT NULL,
  city TEXT NOT NULL,
  temperature DOUBLE PRECISION,
  humidity INTEGER,
  pressure DOUBLE PRECISION,
  wind_speed DOUBLE PRECISION,
  description TEXT,
  api_source TEXT DEFAULT 'openweathermap'
);

-- Create hypertable
SELECT create_hypertable('weather_data', 'time');

-- Stock price data table
CREATE TABLE stock_prices (
  time TIMESTAMPTZ NOT NULL,

```

```

symbol TEXT NOT NULL,
open_price DOUBLE PRECISION,
high_price DOUBLE PRECISION,
low_price DOUBLE PRECISION,
close_price DOUBLE PRECISION,
volume BIGINT,
api_source TEXT DEFAULT 'alphavantage'
);

-- Create hypertable
SELECT create_hypertable('stock_prices', 'time');

-- Sensor data table (for simulated IoT)
CREATE TABLE sensor_readings (
    time TIMESTAMPTZ NOT NULL,
    device_id TEXT NOT NULL,
    sensor_type TEXT NOT NULL,
    value DOUBLE PRECISION,
    unit TEXT,
    location TEXT,
    metadata JSONB
);
;

-- Create hypertable
SELECT create_hypertable('sensor_readings', 'time');

-- Create some indexes for better query performance
CREATE INDEX idx_weather_city_time ON weather_data (city, time DESC);
CREATE INDEX idx_stock_symbol_time ON stock_prices (symbol, time DESC);
CREATE INDEX idx_sensor_device_time ON sensor_readings (device_id, time DESC);

```

Start the Database

```

# Start TimescaleDB
docker compose up -d

# Wait for startup
sleep 10

# Verify it's running
docker compose ps

```

Step 2: API Data Ingestion

Install Python Dependencies

```
# Create virtual environment
python3 -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate

# Install required packages
pip install psycopg2-binary requests pandas matplotlib seaborn python-dotenv schedule
```

Create Environment Configuration

Create `.env` file:

```
# Database connection
DB_HOST=localhost
DB_PORT=5432
DB_NAME=integrations_db
DB_USER=admin
DB_PASSWORD=admin123

# API Keys (Optional - we'll use free endpoints)
OPENWEATHER_API_KEY=your_api_key_here
ALPHAVANTAGE_API_KEY=your_api_key_here
```

Create Weather API Ingestion

Create `weather_ingestion.py`:

```
#!/usr/bin/env python3
import requests
import psycopg2
import json
import time
from datetime import datetime
from dotenv import load_dotenv
import os

load_dotenv()

class WeatherIngestion:
```

```

def __init__(self):
    self.db_config = {
        'host': os.getenv('DB_HOST', 'localhost'),
        'port': os.getenv('DB_PORT', 5432),
        'database': os.getenv('DB_NAME', 'integrations_db'),
        'user': os.getenv('DB_USER', 'admin'),
        'password': os.getenv('DB_PASSWORD', 'admin123')
    }

    # Using free weather API (no key required)
    self.base_url = "https://api.open-meteo.com/v1/forecast"

    # Cities to monitor
    self.cities = {
        'New York': {'lat': 40.7128, 'lon': -74.0060},
        'London': {'lat': 51.5074, 'lon': -0.1278},
        'Tokyo': {'lat': 35.6762, 'lon': 139.6503},
        'Sydney': {'lat': -33.8688, 'lon': 151.2093},
        'Mumbai': {'lat': 19.0760, 'lon': 72.8777}
    }

def connect_db(self):
    """Connect to TimescaleDB"""
    try:
        conn = psycopg2.connect(**self.db_config)
        return conn
    except Exception as e:
        print(f"❌ Database connection failed: {e}")
        return None

def fetch_weather_data(self, city, coordinates):
    """Fetch weather data from Open-Meteo API"""
    try:
        params = {
            'latitude': coordinates['lat'],
            'longitude': coordinates['lon'],
            'current': 'temperature_2m,relative_humidity_2m,surface_pressure,wind_speed_10m',
            'timezone': 'auto'
        }

        response = requests.get(self.base_url, params=params, timeout=10)
        response.raise_for_status()

        data = response.json()
    
```

```

current = data['current']

return {
    'temperature': current.get('temperature_2m'),
    'humidity': current.get('relative_humidity_2m'),
    'pressure': current.get('surface_pressure'),
    'wind_speed': current.get('wind_speed_10m'),
    'description': f"Temperature: {current.get('temperature_2m')}°C",
    'timestamp': datetime.fromisoformat(current['time'].replace('Z', '+00:00'))
}
}

except Exception as e:
    print(f"❌ Failed to fetch weather for {city}: {e}")
    return None

def insert_weather_data(self, city, weather_data):
    """Insert weather data into TimescaleDB"""
    conn = self.connect_db()
    if not conn:
        return False

    try:
        cursor = conn.cursor()

        insert_query = """
            INSERT INTO weather_data (time, city, temperature, humidity, pressure, wind_speed,
            description)
            VALUES (%s, %s, %s, %s, %s, %s)
            """
        cursor.execute(insert_query, (
            weather_data['timestamp'],
            city,
            weather_data['temperature'],
            weather_data['humidity'],
            weather_data['pressure'],
            weather_data['wind_speed'],
            weather_data['description']
        ))

        conn.commit()
        print(f"✅ Inserted weather data for {city}: {weather_data['temperature']}°C")
        return True
    
```

```

except Exception as e:
    print(f"✗ Failed to insert data for {city}: {e}")
    return False
finally:
    conn.close()

def run_ingestion_cycle(self):
    """Run one complete ingestion cycle for all cities"""
    print(f"⌚ Starting weather ingestion at {datetime.now()}")

    success_count = 0
    for city, coordinates in self.cities.items():
        weather_data = self.fetch_weather_data(city, coordinates)

        if weather_data:
            if self.insert_weather_data(city, weather_data):
                success_count += 1
            time.sleep(1) # Be nice to the API

    print(f"📊 Ingestion complete: {success_count}/{len(self.cities)} cities updated")
    return success_count

def run_continuous(self, interval_minutes=15):
    """Run continuous ingestion"""
    print(f"⌚ Starting continuous weather ingestion (every {interval_minutes} minutes)")
    print("Press Ctrl+C to stop")

    try:
        while True:
            self.run_ingestion_cycle()
            print(f"💤 Sleeping for {interval_minutes} minutes...")
            time.sleep(interval_minutes * 60)
    except KeyboardInterrupt:
        print("\n🛑 Stopping weather ingestion")

if __name__ == "__main__":
    ingestion = WeatherIngestion()

    # Run a single cycle first
    print("⌚ Testing single ingestion cycle...")
    ingestion.run_ingestion_cycle()

    # Ask user if they want continuous ingestion

```

```

print("\n" + "*50)
response = input("Run continuous ingestion? (y/n): ").lower()

if response == 'y':
    ingestion.run_continuous(interval_minutes=5) # Every 5 minutes for demo
else:
    print("☑ Single ingestion complete!")

```

Create Simulated IoT Sensor Ingestion

Create `sensor_ingestion.py`:

```

#!/usr/bin/env python3
import psycopg2
import random
import time
import json
from datetime import datetime, timedelta
from dotenv import load_dotenv
import os
import threading

load_dotenv()

class SensorIngestion:
    def __init__(self):
        self.db_config = {
            'host': os.getenv('DB_HOST', 'localhost'),
            'port': os.getenv('DB_PORT', 5432),
            'database': os.getenv('DB_NAME', 'integrations_db'),
            'user': os.getenv('DB_USER', 'admin'),
            'password': os.getenv('DB_PASSWORD', 'admin123')
        }

    # Simulated sensor configurations
    self.sensors = {
        'factory_floor': {
            'devices': ['TEMP_001', 'TEMP_002', 'HUMID_001', 'PRESS_001'],
            'location': 'Factory Floor A',
            'sensors': {
                'TEMP_001': {'type': 'temperature', 'unit': 'celsius', 'range': (18, 26)},
                'TEMP_002': {'type': 'temperature', 'unit': 'celsius', 'range': (20, 28)},
                'HUMID_001': {'type': 'humidity', 'unit': 'percent', 'range': (40, 70)},
            }
        }
    }

```

```

        'PRESS_001': {'type': 'pressure', 'unit': 'hPa', 'range': (1010, 1025)}
    }
},
'warehouse': {
    'devices': ['TEMP_003', 'MOTION_001', 'LIGHT_001'],
    'location': 'Warehouse B',
    'sensors': {
        'TEMP_003': {'type': 'temperature', 'unit': 'celsius', 'range': (15, 25)},
        'MOTION_001': {'type': 'motion', 'unit': 'boolean', 'range': (0, 1)},
        'LIGHT_001': {'type': 'light', 'unit': 'lux', 'range': (100, 800)}
    }
}
}

self.running = False

def connect_db(self):
    """Connect to TimescaleDB"""
    try:
        conn = psycopg2.connect(**self.db_config)
        return conn
    except Exception as e:
        print(f"✗ Database connection failed: {e}")
        return None

def generate_sensor_reading(self, device_id, sensor_config, location):
    """Generate realistic sensor reading"""
    sensor_type = sensor_config['type']
    unit = sensor_config['unit']
    value_range = sensor_config['range']

    if sensor_type == 'motion':
        # Motion sensor: mostly 0, occasionally 1
        value = 1 if random.random() < 0.1 else 0
    elif sensor_type == 'temperature':
        # Temperature with some natural variation
        base_temp = sum(value_range) / 2
        variation = (value_range[1] - value_range[0]) * 0.3
        value = base_temp + random.uniform(-variation, variation)
        value = round(value, 2)
    elif sensor_type == 'humidity':
        # Humidity with gradual changes
        value = random.uniform(*value_range)
        value = round(value, 1)

```

```

elif sensor_type == 'pressure':
    # Pressure with small variations
    base_pressure = sum(value_range) / 2
    variation = (value_range[1] - value_range[0]) * 0.2
    value = base_pressure + random.uniform(-variation, variation)
    value = round(value, 2)
else:
    # Default: random value in range
    value = round(random.uniform(*value_range), 2)

# Add some metadata
metadata = {
    'battery_level': random.randint(75, 100),
    'signal_strength': random.randint(-80, -30),
    'firmware_version': '1.2.3'
}

return {
    'device_id': device_id,
    'sensor_type': sensor_type,
    'value': value,
    'unit': unit,
    'location': location,
    'metadata': json.dumps(metadata),
    'timestamp': datetime.now()
}

def insert_sensor_reading(self, reading):
    """Insert sensor reading into TimescaleDB"""
    conn = self.connect_db()
    if not conn:
        return False

    try:
        cursor = conn.cursor()

        insert_query = """
            INSERT INTO sensor_readings (time, device_id, sensor_type, value, unit, location,
            metadata)
            VALUES (%s, %s, %s, %s, %s, %s)
        """

        cursor.execute(insert_query, (
            reading['timestamp'],

```

```

        reading['device_id'],
        reading['sensor_type'],
        reading['value'],
        reading['unit'],
        reading['location'],
        reading['metadata']
    ))
    conn.commit()
    print(f"⌚ {reading['device_id']}: {reading['value']} {reading['unit']}")
    return True
except Exception as e:
    print(f"✗ Failed to insert sensor reading: {e}")
    return False
finally:
    conn.close()

def simulate_location_sensors(self, location_name, location_config):
    """Simulate all sensors for a specific location"""
    while self.running:
        for device_id, sensor_config in location_config["sensors"].items():
            reading = self.generate_sensor_reading(
                device_id,
                sensor_config,
                location_config["location"]
            )
            self.insert_sensor_reading(reading)
            time.sleep(0.5) # Small delay between readings

        time.sleep(5) # Delay between cycles

def start_simulation(self, duration_minutes=10):
    """Start sensor simulation for all locations"""
    print(f"⌚ Starting IoT sensor simulation for {duration_minutes} minutes")
    print("⌚ Sensors active:")

    for location, config in self.sensors.items():
        print(f"⌚ {config['location']}: {''.join(config['devices'])}")

    self.running = True
    threads = []

```

```

# Start a thread for each location
for location_name, location_config in self.sensors.items():
    thread = threading.Thread(
        target=self.simulate_location_sensors,
        args=(location_name, location_config),
        daemon=True
    )
    thread.start()
    threads.append(thread)

try:
    # Run for specified duration
    time.sleep(duration_minutes * 60)
except KeyboardInterrupt:
    print("\n🔴 Stopping sensor simulation...")
finally:
    self.running = False
    print("✅ Sensor simulation stopped")

if __name__ == "__main__":
    simulation = SensorIngestion()

    print("📝 Testing sensor ingestion...")
    simulation.start_simulation(duration_minutes=2) # Run for 2 minutes

```

Run the Data Ingestion

```

# Terminal 1: Start weather ingestion
python weather_ingestion.py

# Terminal 2: Start sensor simulation (run this after weather starts)
python sensor_ingestion.py

```

Step 3: Data Export and Analysis with Pandas

Create Data Analysis Module

Create `data_analysis.py`:

```

#!/usr/bin/env python3
import pandas as pd
import psycopg2

```

```

import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime, timedelta
import numpy as np
from dotenv import load_dotenv
import os

load_dotenv()

class TimescaleAnalyzer:
    def __init__(self):
        self.db_config = {
            'host': os.getenv('DB_HOST', 'localhost'),
            'port': os.getenv('DB_PORT', 5432),
            'database': os.getenv('DB_NAME', 'integrations_db'),
            'user': os.getenv('DB_USER', 'admin'),
            'password': os.getenv('DB_PASSWORD', 'admin123')
        }

    # Set up plotting style
    plt.style.use('seaborn-v0_8')
    sns.set_palette("husl")

    def connect_db(self):
        """Create database connection"""
        try:
            conn = psycopg2.connect(**self.db_config)
            return conn
        except Exception as e:
            print(f"X Database connection failed: {e}")
            return None

    def query_to_dataframe(self, query, params=None):
        """Execute query and return pandas DataFrame"""
        conn = self.connect_db()
        if not conn:
            return None

        try:
            df = pd.read_sql_query(query, conn, params=params)
            return df
        except Exception as e:
            print(f"X Query failed: {e}")
            return None

```

```

finally:
    conn.close()

def analyze_weather_data(self):
    """Analyze weather data and create visualizations"""
    print("🌩️ Analyzing weather data...")

    # Query recent weather data
    query = """
        SELECT
            time,
            city,
            temperature,
            humidity,
            pressure,
            wind_speed
        FROM weather_data
        WHERE time >= NOW() - INTERVAL '24 hours'
        ORDER BY time DESC, city
    """

    df = self.query_to_dataframe(query)

    if df is None or df.empty:
        print("❌ No weather data found")
        return

    print(f"📊 Found {len(df)} weather records")
    print(f"🏙️ Cities: {', '.join(df['city'].unique())}")

    # Convert time to datetime
    df['time'] = pd.to_datetime(df['time'])

    # Create visualizations
    fig, axes = plt.subplots(2, 2, figsize=(15, 12))
    fig.suptitle('Weather Data Analysis', fontsize=16, fontweight='bold')

    # Temperature by city over time
    for city in df['city'].unique():
        city_data = df[df['city'] == city]
        axes[0, 0].plot(city_data['time'], city_data['temperature'],
                        marker='o', label=city, linewidth=2)

    axes[0, 0].set_title('Temperature Over Time by City')

```

```

axes[0, 0].set_xlabel('Time')
axes[0, 0].set_ylabel('Temperature (°C)')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)
axes[0, 0].tick_params(axis='x', rotation=45)

# Current temperature comparison
latest_data = df.groupby('city').last()
bars = axes[0, 1].bar(latest_data.index, latest_data['temperature'])
axes[0, 1].set_title('Current Temperature by City')
axes[0, 1].set_ylabel('Temperature (°C)')
axes[0, 1].tick_params(axis='x', rotation=45)

# Add value labels on bars
for bar in bars:
    height = bar.get_height()
    axes[0, 1].text(bar.get_x() + bar.get_width()/2., height,
                    f'{height:.1f}°C',
                    ha='center', va='bottom')

# Humidity vs Temperature scatter
colors = plt.cm.viridis(np.linspace(0, 1, len(df['city'].unique())))
for i, city in enumerate(df['city'].unique()):
    city_data = df[df['city'] == city]
    axes[1, 0].scatter(city_data['temperature'], city_data['humidity'],
                       label=city, alpha=0.7, s=60, color=colors[i])

    axes[1, 0].set_title('Humidity vs Temperature')
    axes[1, 0].set_xlabel('Temperature (°C)')
    axes[1, 0].set_ylabel('Humidity (%)')
    axes[1, 0].legend()
    axes[1, 0].grid(True, alpha=0.3)

# Weather summary statistics
summary_stats = df.groupby('city').agg({
    'temperature': ['mean', 'min', 'max'],
    'humidity': 'mean',
    'wind_speed': 'mean'
}).round(2)

# Create a heatmap of average values
heatmap_data = df.groupby('city')[['temperature', 'humidity', 'wind_speed']].mean()
im = axes[1, 1].imshow(heatmap_data.T, cmap='YlOrRd', aspect='auto')
axes[1, 1].set_title('Average Weather Metrics Heatmap')

```

```

axes[1, 1].set_xticks(range(len(heatmap_data.index)))
axes[1, 1].set_xticklabels(heatmap_data.index, rotation=45)
axes[1, 1].set_yticks(range(len(heatmap_data.columns)))
axes[1, 1].set_yticklabels(heatmap_data.columns)

# Add colorbar
plt.colorbar(im, ax=axes[1, 1])

# Add text annotations
for i in range(len(heatmap_data.index)):
    for j in range(len(heatmap_data.columns)):
        text = axes[1, 1].text(i, j, f'{heatmap_data.iloc[i, j]:.1f}',
                               ha="center", va="center", color="black")

plt.tight_layout()
plt.savefig('weather_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

# Print summary statistics
print("\nweathermap Weather Summary Statistics:")
print(summary_stats)

return df

def analyze_sensor_data(self):
    """Analyze IoT sensor data"""
    print("\nweather Analyzing sensor data...")

    # Query recent sensor data
    query = """
        SELECT
            time,
            device_id,
            sensor_type,
            value,
            unit,
            location,
            metadata
        FROM sensor_readings
        WHERE time >= NOW() - INTERVAL '2 hours'
        ORDER BY time DESC
    """

    df = self.query_to_dataframe(query)

```

```

if df is None or df.empty:
    print("X No sensor data found")
    return

print(f"Found {len(df)} sensor readings")
print(f"Locations: {', '.join(df['location'].unique())}")

# Convert time to datetime
df['time'] = pd.to_datetime(df['time'])

# Create sensor-specific visualizations
fig, axes = plt.subplots(2, 2, figsize=(15, 12))
fig.suptitle('IoT Sensor Data Analysis', fontsize=16, fontweight='bold')

# Temperature sensors over time
temp_data = df[df['sensor_type'] == 'temperature']
if not temp_data.empty:
    for device in temp_data['device_id'].unique():
        device_data = temp_data[temp_data['device_id'] == device]
        axes[0, 0].plot(device_data['time'], device_data['value'],
                         marker='o', label=device, linewidth=2)

    axes[0, 0].set_title('Temperature Sensors Over Time')
    axes[0, 0].set_xlabel('Time')
    axes[0, 0].set_ylabel('Temperature (°C)')
    axes[0, 0].legend()
    axes[0, 0].grid(True, alpha=0.3)
    axes[0, 0].tick_params(axis='x', rotation=45)

# Sensor type distribution
sensor_counts = df['sensor_type'].value_counts()
bars = axes[0, 1].bar(sensor_counts.index, sensor_counts.values)
axes[0, 1].set_title('Sensor Reading Count by Type')
axes[0, 1].set_ylabel('Number of Readings')
axes[0, 1].tick_params(axis='x', rotation=45)

# Add value labels on bars
for bar in bars:
    height = bar.get_height()
    axes[0, 1].text(bar.get_x() + bar.get_width()/2., height,
                    f'{int(height)}',
                    ha='center', va='bottom')

```

```

# Humidity sensor data (if available)
humidity_data = df[df['sensor_type'] == 'humidity']
if not humidity_data.empty:
    axes[1, 0].plot(humidity_data['time'], humidity_data['value'],
                    'g-o', linewidth=2)
    axes[1, 0].set_title('Humidity Sensor Readings')
    axes[1, 0].set_xlabel('Time')
    axes[1, 0].set_ylabel('Humidity (%)')
    axes[1, 0].grid(True, alpha=0.3)
    axes[1, 0].tick_params(axis='x', rotation=45)
else:
    axes[1, 0].text(0.5, 0.5, 'No humidity data available',
                    ha='center', va='center', transform=axes[1, 0].transAxes)

# Latest sensor values by location
latest_readings = df.loc[df.groupby(['location', 'sensor_type'])['time'].idxmax()]
pivot_data = latest_readings.pivot_table(
    index='location',
    columns='sensor_type',
    values='value',
    aggfunc='mean'
)
if not pivot_data.empty:
    sns.heatmap(pivot_data, annot=True, fmt='.1f',
                cmap='viridis', ax=axes[1, 1])
    axes[1, 1].set_title('Latest Sensor Values by Location')
else:
    axes[1, 1].text(0.5, 0.5, 'Insufficient data for heatmap',
                    ha='center', va='center', transform=axes[1, 1].transAxes)

plt.tight_layout()
plt.savefig('sensor_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

# Print summary statistics
print("\n\x02 Sensor Summary Statistics:")
summary = df.groupby(['location', 'sensor_type']).agg({
    'value': ['count', 'mean', 'min', 'max']
}).round(2)
print(summary)

return df

```

```

def create_combined_dashboard(self):
    """Create a combined dashboard with multiple data sources"""
    print("\n📊 Creating combined data dashboard...")

    # Query data from multiple tables
    queries = {
        'weather': """
            SELECT
                time_bucket('1 hour', time) as bucket,
                city,
                AVG(temperature) as avg_temp,
                AVG(humidity) as avg_humidity
            FROM weather_data
            WHERE time >= NOW() - INTERVAL '24 hours'
            GROUP BY bucket, city
            ORDER BY bucket DESC
        """,
        'sensors': """
            SELECT
                time_bucket('15 minutes', time) as bucket,
                sensor_type,
                location,
                AVG(value) as avg_value,
                COUNT(*) as reading_count
            FROM sensor_readings
            WHERE time >= NOW() - INTERVAL '2 hours'
            GROUP BY bucket, sensor_type, location
            ORDER BY bucket DESC
        """
    }

    data = {}
    for name, query in queries.items():
        df = self.query_to_dataframe(query)
        if df is not None and not df.empty:
            df['bucket'] = pd.to_datetime(df['bucket'])
            data[name] = df
        else:
            print(f"⚠️ No data found for {name}")

    if not data:
        print("❌ No data available for dashboard")
        return

```

```

# Create dashboard
fig = plt.figure(figsize=(16, 10))
fig.suptitle('TimescaleDB Integration Dashboard', fontsize=18, fontweight='bold')

# Create grid layout
gs = fig.add_gridspec(3, 3, hspace=0.3, wspace=0.3)

# Weather data visualizations
if 'weather' in data:
    weather_df = data['weather']

    # Temperature trends
    ax1 = fig.add_subplot(gs[0, :2])
    for city in weather_df['city'].unique():
        city_data = weather_df[weather_df['city'] == city]
        ax1.plot(city_data['bucket'], city_data['avg_temp'],
                 marker='o', label=city, linewidth=2)
    ax1.set_title('Weather Temperature Trends (24h)')
    ax1.set_ylabel('Temperature (°C)')
    ax1.legend()
    ax1.grid(True, alpha=0.3)
    ax1.tick_params(axis='x', rotation=45)

    # Current weather summary
    ax2 = fig.add_subplot(gs[0, 2])
    latest_weather = weather_df.groupby('city').last()
    bars = ax2.bar(range(len(latest_weather)), latest_weather['avg_temp'])
    ax2.set_title('Latest Temperatures')
    ax2.set_xticks(range(len(latest_weather)))
    ax2.set_xticklabels(latest_weather.index, rotation=45)
    ax2.set_ylabel('Temperature (°C)')

    # Add value labels
    for i, bar in enumerate(bars):
        height = bar.get_height()
        ax2.text(bar.get_x() + bar.get_width()/2., height,
                 f'{height:.1f}°C',
                 ha='center', va='bottom')

# Sensor data visualizations
if 'sensors' in data:
    sensors_df = data['sensors']

    # Temperature sensor trends

```

```

ax3 = fig.add_subplot(gs[1, :2])
temp_sensors = sensors_df[sensors_df['sensor_type'] == 'temperature']
if not temp_sensors.empty:
    for location in temp_sensors['location'].unique():
        loc_data = temp_sensors[temp_sensors['location'] == location]
        ax3.plot(loc_data['bucket'], loc_data['avg_value'],
                  marker='s', label=location, linewidth=2)
    ax3.set_title('IoT Temperature Sensors (2h)')
    ax3.set_ylabel('Temperature (°C)')
    ax3.legend()
    ax3.grid(True, alpha=0.3)
    ax3.tick_params(axis='x', rotation=45)
else:
    ax3.text(0.5, 0.5, 'No temperature sensor data',
             ha='center', va='center', transform=ax3.transAxes)

# Sensor activity summary
ax4 = fig.add_subplot(gs[1, 2])
activity = sensors_df.groupby('sensor_type')['reading_count'].sum()
wedges, texts, autotexts = ax4.pie(activity.values, labels=activity.index,
                                     autopct='%1.1f%%', startangle=90)
ax4.set_title('Sensor Activity Distribution')

# Combined metrics
ax5 = fig.add_subplot(gs[2, :])

# Create a timeline showing data ingestion rates
timeline_data = []

if 'weather' in data:
    weather_timeline = data['weather'].groupby('bucket').size().reset_index()
    weather_timeline['source'] = 'Weather API'
    weather_timeline.columns = ['time', 'count', 'source']
    timeline_data.append(weather_timeline)

if 'sensors' in data:
    sensor_timeline = data['sensors'].groupby('bucket')['reading_count'].sum().reset_index()
    sensor_timeline['source'] = 'IoT Sensors'
    sensor_timeline.columns = ['time', 'count', 'source']
    timeline_data.append(sensor_timeline)

if timeline_data:
    combined_timeline = pd.concat(timeline_data, ignore_index=True)

```

```

for source in combined_timeline['source'].unique():
    source_data = combined_timeline[combined_timeline['source'] == source]
    ax5.plot(source_data['time'], source_data['count'],
              marker='o', label=source, linewidth=2)

ax5.set_title('Data Ingestion Timeline')
ax5.set_xlabel('Time')
ax5.set_ylabel('Records Count')
ax5.legend()
ax5.grid(True, alpha=0.3)
ax5.tick_params(axis='x', rotation=45)
else:
    ax5.text(0.5, 0.5, 'No timeline data available',
             ha='center', va='center', transform=ax5.transAxes)

plt.savefig('combined_dashboard.png', dpi=300, bbox_inches='tight')
plt.show()

# Print data summary
print("\n📋 Data Integration Summary:")
if 'weather' in data:
    print(f" 🌡 Weather: {len(data['weather'])} hourly averages from
{data['weather']['city'].nunique()} cities")
    if 'sensors' in data:
        print(f" 💡 Sensors: {len(data['sensors'])} readings from
{data['sensors']['location'].nunique()} locations")

return data

def export_data_samples(self):
    """Export sample data to various formats for external use"""
    print("\n💾 Exporting data samples...")

    # Export weather data
    weather_query = """
        SELECT * FROM weather_data
        WHERE time >= NOW() - INTERVAL '24 hours'
        ORDER BY time DESC
    """

    weather_df = self.query_to_dataframe(weather_query)
    if weather_df is not None and not weather_df.empty:
        weather_df.to_csv('weather_export.csv', index=False)
        weather_df.to_json('weather_export.json', orient='records', date_format='iso')

```

```

print(f" ✅ Weather data exported: {len(weather_df)} records")

# Export sensor data
sensor_query = """
    SELECT * FROM sensor_readings
    WHERE time >= NOW() - INTERVAL '2 hours'
    ORDER BY time DESC
"""

sensor_df = self.query_to_dataframe(sensor_query)
if sensor_df is not None and not sensor_df.empty:
    sensor_df.to_csv('sensor_export.csv', index=False)
    sensor_df.to_json('sensor_export.json', orient='records', date_format='iso')
    print(f" ✅ Sensor data exported: {len(sensor_df)} records")

# Create summary report
if weather_df is not None and sensor_df is not None:
    summary_report = {
        'export_timestamp': datetime.now().isoformat(),
        'weather_records': len(weather_df),
        'sensor_records': len(sensor_df),
        'weather_cities': weather_df['city'].nunique() if not weather_df.empty else 0,
        'sensor_locations': sensor_df['location'].nunique() if not sensor_df.empty else 0,
        'time_range': {
            'start': min(weather_df['time'].min() if not weather_df.empty else datetime.now(),
                         sensor_df['time'].min() if not sensor_df.empty else
                         datetime.now()).isoformat(),
            'end': max(weather_df['time'].max() if not weather_df.empty else datetime.now(),
                      sensor_df['time'].max() if not sensor_df.empty else
                      datetime.now()).isoformat()
        }
    }

    import json
    with open('data_export_summary.json', 'w') as f:
        json.dump(summary_report, f, indent=2)

    print(f" ✅ Summary report created")

print("💾 Export complete! Files created:")
print(" - weather_export.csv/json")
print(" - sensor_export.csv/json")
print(" - data_export_summary.json")

```

```

if __name__ == "__main__":
    analyzer = TimescaleAnalyzer()

    print("⌚ Starting TimescaleDB Data Analysis")
    print("=" * 50)

    # Run all analyses
    try:
        weather_data = analyzer.analyze_weather_data()
        sensor_data = analyzer.analyze_sensor_data()
        combined_data = analyzer.create_combined_dashboard()
        analyzer.export_data_samples()

        print("\n✅ Analysis complete! Check the generated PNG files and exported data.")

    except Exception as e:
        print(f"❌ Analysis failed: {e}")

```

Create Simple Pipeline Runner

Create `run_pipeline.py`:

```

#!/usr/bin/env python3
import subprocess
import time
import sys
import os
from datetime import datetime

def run_command(command, description):
    """Run a command and handle errors"""
    print(f"\n⌚ {description}")
    print(f"💻 Running: {command}")

    try:
        result = subprocess.run(command, shell=True, check=True,
                               capture_output=True, text=True)
        print(f"✅ {description} completed successfully")
        if result.stdout:
            print(f"📄 Output: {result.stdout.strip()}")
        return True
    except subprocess.CalledProcessError as e:

```

```

print(f"✗ {description} failed")
print(f"Error: {e.stderr.strip() if e.stderr else str(e)}")
return False

def check_dependencies():
    """Check if all required packages are installed"""
    print("⌚ Checking dependencies...")

    required_packages = ['psycopg2', 'pandas', 'matplotlib', 'seaborn', 'requests']
    missing_packages = []

    for package in required_packages:
        try:
            __import__(package)
            print(f"✓ {package}")
        except ImportError:
            print(f"✗ {package} - MISSING")
            missing_packages.append(package)

    if missing_packages:
        print("\n⚠ Missing packages: {}, '.join(missing_packages))")
        print("💡 Run: pip install psycopg2-binary pandas matplotlib seaborn requests python-dotenv")
        return False

    print("✓ All dependencies satisfied")
    return True

def main():
    print("⌚ TimescaleDB Integration Pipeline Runner")
    print("=" * 50)

    # Check dependencies
    if not check_dependencies():
        sys.exit(1)

    # Check if database is running
    print("\n⌚ Checking database connection...")
    if not run_command("docker compose ps timescaledb", "Database status check"):
        print("💡 Start the database with: docker compose up -d")
        sys.exit(1)

```

```

print("\n📋 Pipeline Options:")
print("1. Run complete pipeline (ingestion + analysis)")
print("2. Run data ingestion only")
print("3. Run data analysis only")
print("4. Export data samples")
print("5. Check data status")

choice = input("\nSelect option (1-5): ").strip()

if choice == "1":
    # Complete pipeline
    print("\n⌚️ Running complete data pipeline...")

    # Start weather ingestion in background
    print("☁️ Starting weather data ingestion...")
    weather_process = subprocess.Popen([sys.executable, 'weather_ingestion.py'])

    # Wait a bit for weather data to start
    time.sleep(10)

    # Start sensor simulation in background
    print("📡 Starting sensor simulation...")
    sensor_process = subprocess.Popen([sys.executable, 'sensor_ingestion.py'])

    # Wait for data collection
    print("🕒 Collecting data for 3 minutes...")
    time.sleep(180) # 3 minutes

    # Stop ingestion processes
    print("🛑 Stopping data ingestion...")
    weather_process.terminate()
    sensor_process.terminate()

    # Wait a moment for cleanup
    time.sleep(5)

    # Run analysis
    print("📊 Running data analysis...")
    run_command(f"{sys.executable} data_analysis.py", "Data analysis")

elif choice == "2":
    # Ingestion only
    print("\n📋 Starting data ingestion...")

```

```

print("☁️ Weather ingestion will run for 2 minutes")
print("⚡ Sensor simulation will run for 2 minutes")
print("Press Ctrl+C to stop early")

try:
    # Start both ingestion processes
    weather_process = subprocess.Popen([sys.executable, 'weather_ingestion.py'])
    time.sleep(5)
    sensor_process = subprocess.Popen([sys.executable, 'sensor_ingestion.py'])

    # Wait for data collection
    time.sleep(120) # 2 minutes

    # Stop processes
    weather_process.terminate()
    sensor_process.terminate()
    print("✅ Data ingestion completed")

except KeyboardInterrupt:
    print("\n🛑 Stopping ingestion processes...")
    try:
        weather_process.terminate()
        sensor_process.terminate()
    except:
        pass

elif choice == "3":
    # Analysis only
    run_command(f"{sys.executable} data_analysis.py", "Data analysis")

elif choice == "4":
    # Export only
    print("\nEXPORTING")
    run_command(f"{sys.executable} -c \"from data_analysis import TimescaleAnalyzer;
TimescaleAnalyzer().export_data_samples()\"", "Data export")

elif choice == "5":
    # Status check
    run_command(f"{sys.executable} -c \"from data_analysis import TimescaleAnalyzer;
analyzer = TimescaleAnalyzer(); analyzer.query_to_dataframe('SELECT COUNT(*) as
weather_records FROM weather_data WHERE time >= NOW() - INTERVAL \'24 hours\';');
analyzer.query_to_dataframe('SELECT COUNT(*) as sensor_records FROM sensor_readings
WHERE time >= NOW() - INTERVAL \'2 hours\';')\"", "Data status check")

```

```

else:
    print("X Invalid option selected")
    sys.exit(1)

print(f"\n

```

Manual Step-by-Step Execution

If you prefer to run each component manually:

```
# Terminal 1: Start weather ingestion (let it run for a few minutes)
python weather_ingestion.py
```

```
# Terminal 2: Start sensor simulation (let it run for a few minutes)
python sensor_ingestion.py
```

```
# Terminal 3: After collecting data, run analysis
python data_analysis.py
```

Step 5: Understanding the Integration Patterns

Integration Architecture

External APIs → TimescaleDB → Python/Pandas → Visualizations



Weather API	Hypertables	DataFrames	Matplotlib
IoT Sensors	Time-series	Analysis	Dashboards
REST APIs	Compression	Aggregation	Exports
File Imports	SQL Queries	Statistics	CSV/JSON

Key Integration Benefits

TimescaleDB Advantages:

- **Time-series optimization:** Automatic partitioning and compression
- **SQL compatibility:** Use familiar SQL for complex queries
- **Scalability:** Handle millions of time-series data points
- **Real-time ingestion:** Insert data as it arrives

Python/Pandas Benefits:

- **Rich analysis:** Statistical functions and data manipulation
- **Visualization:** Matplotlib, Seaborn for charts and graphs
- **Machine learning:** Easy integration with scikit-learn, TensorFlow
- **Export flexibility:** CSV, JSON, Excel, and more

Common Integration Patterns

Ingestion Patterns:

1. **Polling:** Regular API calls (weather example)
2. **Streaming:** Real-time data feeds (sensor simulation)
3. **Batch:** Scheduled bulk uploads
4. **Event-driven:** Triggered by external events

Export Patterns:

1. **Scheduled reports:** Daily/weekly summaries
2. **Real-time dashboards:** Live data visualization
3. **File exports:** CSV, JSON for external analysis
4. **Python analysis:** Feed data science workflows

Step 6: Production Considerations

```
# Connect to TimescaleDB using psql
docker exec -it timescale-integrations psql -U admin -d integrations_db
```

Performance Optimization

```
-- Create BRIN index
CREATE INDEX idx_weather_city_time_brin
ON weather_data USING BRIN (city, time);

-- Create continuous aggregate for real-time analytics
CREATE MATERIALIZED VIEW weather_hourly
WITH (timescaledb.continuous) AS
SELECT
    time_bucket('1 hour', time) AS bucket,
    city,
    AVG(temperature) as avg_temp,
    MIN(temperature) as min_temp,
    MAX(temperature) as max_temp,
    AVG(humidity) as avg_humidity
FROM weather_data
GROUP BY bucket, city;

-- Refresh policy for real-time updates
SELECT add_continuous_aggregate_policy('weather_hourly',
    start_offset => INTERVAL '1 day',
    end_offset => INTERVAL '1 hour',
    schedule_interval => INTERVAL '1 hour');
```

Error Handling and Monitoring

Create `monitoring.py`:

```
#!/usr/bin/env python3
import psycopg2
from datetime import datetime, timedelta
import json

def check_data_health():
    """Check data pipeline health"""
    checks = {
        'weather_data_recent': """
            SELECT COUNT(*) as count
            FROM weather_data
            WHERE time >= NOW() - INTERVAL '1 hour'
        """,
        'sensor_data_recent': """
            SELECT COUNT(*) as count
        """}
```

```

    FROM sensor_readings
    WHERE time >= NOW() - INTERVAL '15 minutes'
    """,
    'total_records': """
        SELECT
            (SELECT COUNT(*) FROM weather_data) as weather_count,
            (SELECT COUNT(*) FROM sensor_readings) as sensor_count
        """
    }

# Implementation would check each query and alert if thresholds not met
print("🔍 Data pipeline health check would run here")
print("✅ All checks passed (demo)")

if __name__ == "__main__":
    check_data_health()

```

Security Best Practices

```

# Example: Environment-based configuration
import os
from urllib.parse import quote_plus

# Secure database connection
DB_CONFIG = {
    'host': os.getenv('DB_HOST', 'localhost'),
    'port': int(os.getenv('DB_PORT', 5432)),
    'database': os.getenv('DB_NAME'),
    'user': os.getenv('DB_USER'),
    'password': quote_plus(os.getenv('DB_PASSWORD')),
    'sslmode': os.getenv('DB_SSL_MODE', 'require')
}

# API key management
API_KEYS = {
    'weather': os.getenv('WEATHER_API_KEY'),
    'stock': os.getenv('STOCK_API_KEY')
}

```

Step 7: Cleanup

Stop All Processes

```
# If running continuous ingestion, stop with Ctrl+C
```

```
# Stop the database
```

```
docker compose down -v
```

```
# Deactivate virtual environment
```

```
deactivate
```

```
# Remove project directory (optional)
```

```
cd ..
```

```
rm -rf timescale-integrations
```

Verify Cleanup

```
# Check no containers remain
```

```
docker ps -a | grep timescale
```

```
# Check system usage
```

```
docker system df
```

Summary: What You Learned

Core Concepts Mastered:

- **Data ingestion patterns:** REST API polling, IoT simulation, real-time streaming
- **TimescaleDB integration:** Hypertables, time-series optimization, SQL queries
- **Python data pipeline:** Pandas for analysis, Matplotlib for visualization
- **Export mechanisms:** CSV, JSON, and programmatic data access

Hands-On Skills:

- Built weather data ingestion from public APIs
- Simulated IoT sensor data streams
- Created comprehensive data analysis with Pandas
- Generated professional visualizations with Matplotlib
- Exported data in multiple formats for external use

Production-Ready Patterns:

1. **Modular architecture:** Separate ingestion, analysis, and visualization
2. **Error handling:** Robust connection management and graceful failures
3. **Configuration management:** Environment variables and secure secrets
4. **Monitoring:** Data quality checks and pipeline health monitoring

Key Integration Benefits:

- **Real-time insights:** From raw data to visualizations in minutes
- **Scalable architecture:** Handle growing data volumes efficiently
- **Flexible exports:** Data available in multiple formats for different tools
- **Time-series power:** Leverage TimescaleDB's optimization for time-based data

Next Steps for Your Projects:

- Start with simple API integrations using the patterns shown
- Build incrementally: ingestion → storage → analysis → visualization
- Monitor data quality and pipeline performance from day one
- Scale up gradually as your data volume and complexity grows

You now have a complete understanding of building data integration pipelines with TimescaleDB and can create production-ready systems for your own time-series data projects!